Solderbench: An open-source synthetic benchmark for PCB solder-joint defect inspection algorithm development

Qinmin Wang[a,1], Zixin Wang[b,2], Hongyuan Wang[a,3], Zongzhong Wang[c,*]

[a]*School of Computer Science and Artificial Intelligence, Guangdong University of Science and Technology, Dongguan, China*
[b]*School of Electronic and Communication Engineering, Tianjin Normal University, Tianjin, China*
[c]*School of Electromechanical Engineering, Guangdong University of Technology, Guangzhou, China*

Abstract

Automated optical inspection (AOI) of solder joints is a critical quality-control step in printed circuit board (PCB) assembly, and three-dimensional (3D) point-cloud inspection is increasingly adopted because solder quality is essentially a volumetric property. However, publicly available 3D solder-joint datasets are essentially nonexistent: industrial scan data are proprietary, so new inspection algorithms are difficult to compare fairly. We present Solderbench, an open-source, pure-Python synthetic benchmark for developing and evaluating 3D solder-joint defect inspection algorithms. Solderbench provides (i) a parametric point-cloud generator that simulates five joint classes (good, insufficient, excessive, bridging, and collapse) with more than fifteen controllable parameters and seed-level, point-wise reproducibility; (ii) a seven-dimensional geometric feature extractor; (iii) six baseline classifiers, including a re-implementation of the industrial volume-threshold rule S3DA; and (iv) automated comparison and ablation runners that export CSV, LaTeX, and figure artifacts. On a 1250-sample benchmark (seed 42), a random forest reaches 0.9440 accuracy and 0.9439 macro-F_1, and feature ablation identifies height variance as the most informative descriptor. Synthetic results serve as a controlled, reproducible proxy for algorithm development, yet final validation must be performed using real-world production scan data.

Keywords: solder joint inspection; automated optical inspection; 3D point cloud; synthetic benchmark; defect classification; reproducible evaluation; PCB; synthetic dataset

1. Required metadata

2. Motivation and significance

Surface-mount technology (SMT) dominates modern printed-circuit-board (PCB) assembly, and solder-joint quality largely determines product reliability. AOI is the standard quality-control step on SMT lines [1, 2, 3]; classical 2D AOI and deep-learning 2D detectors [4] on public datasets such as DeepPCB [5] are mature, yet inspection is increasingly moving to 3D because solder quality is fundamentally volumetric [6, 7]. General 3D deep learning (PointNet/PointNet++ [8, 9], Open3D [10]) has advanced, but 3D solder-joint research remains hampered by the absence of public data: existing public PCB datasets are almost entirely 2D [5], and the few 3D studies rely on private, non-public in-house point-cloud datasets [6, 7], so algorithms cannot be compared fairly or reproduced. Solderbench closes this gap synthetically, simulating solder-joint point clouds from a parametric geometric model whose parameters (nominal joint size, noise, defect strength, bridging distance, collapse depth) are physically interpretable and user-controllable; synthetic generation yields costless labeled data of arbitrary size, deliberate control of class difficulty and sensor noise, and seed-level reproducibility.

*Corresponding author at School of Electromechanical Engineering, Guangdong University of Technology, Guangzhou, China. E-mail: wangzz@gdut.edu.cn
[1]E-mail: wangqinmin@gdust.edu.cn
[2]E-mail: 396514839@qq.com
[3]E-mail: wanghongyuan@gdust.edu.cn

Table 1: Code metadata (mandatory).

Nr.	Code metadata description	Value
C1	Current code version	v0.1.0
C2	Permanent link to code/repository used for this code version	https://github.com/OWNER/solderbench
C3	Permanent link to reproducible capsule	10.5281/zenodo.XXXXXXX
C4	Legal code license	MIT License
C5	Code versioning system used	git
C6	Software code languages, tools, and services used	Python ($\geq$ 3.8)
C7	Compilation requirements, operating environments and dependencies	NumPy, SciPy, pandas, scikit-learn, and Matplotlib; OS-independent (no compilation, no GPU required)
C8	Link to developer documentation/manual	https://github.com/OWNER/solderbench#readme
C9	Support email for questions	wangzz@gdut.edu.cn

3. Software description

3.1. Software architecture

Solderbench is a pure-Python package (about 1,000 source lines) with five feed-forward modules behind one public API: `generator.py` (parametric simulator), `features.py` (seven descriptors), `baselines.py` (six classifiers), `evaluate.py` (`run_comparison`/`run_ablation` runners exporting CSV, Booktabs-style LaTeX tables, and figures), and `data.py` (`load_dataset()`, the real-camera integration point). Stages exchange NumPy arrays and pandas DataFrames, so any stage is replaceable; synthetic clouds swap for segmented real scans at `load_dataset()` without downstream changes. The package ships a 15-test suite (shape and class-balance checks, same-seed reproducibility, feature validity, a Welch t-test that insufficient-solder volumes are significantly smaller, and fit-and-predict smoke tests), a `pyproject.toml`, and a `codemeta.json`; dependencies are only NumPy, SciPy, pandas, scikit-learn, and Matplotlib [14, 15, 16, 17].

3.2. Software functionalities

Generation. SolderJointGenerator simulates five SMT defect classes [6, 7]: good, insufficient (shrunk ellipsoid, 0.55–0.85 of nominal), excessive (inflated ellipsoid, mean inflation 1.3, plus a bump), bridging (small center-to-center distance plus a strip), and collapse (random top dent). Over fifteen parameters control n_{points}, noise, and the shrinkage, inflation, bridging, and collapse ranges; class ranges deliberately overlap so no classifier is perfect, and difficulty is tuned by widening or narrowing them.

Features. extract_features returns seven descriptors (Table 2): `vol` (voxel volume), `max_h`, `h_var`, `area`, `vol_area_ratio`, `convex_vol_ratio` (irregularity), and `min_dist` (bridging exposure). **Baselines and**

Table 2: The seven geometric descriptors returned by `extract_features` and their physical meaning.

Feature	Definition	Physical meaning
`vol`	voxel count $\times$ (voxel size)^3	total solder volume
`max_h`	max z over points	joint height
`h_var`	variance of point heights	height-distribution spread (dent / bump)
`area`	convex-hull surface area	joint surface extent
`vol_area_ratio`	vol / area	compactness
`convex_vol_ratio`	vol / ConvexHull.volume	shape irregularity (voids / dents)
`min_dist`	min distance to neighbouring joint	bridging exposure

runners. get_baselines returns S3DA — a re-implementation of the industrial volume-threshold rule [6], with configurable thresholds (default values hard-coded in source code) — plus logistic regression, an RBF-SVM [12], kNN, gradient boosting, and a random forest [11]; all are scikit-learn estimators [13] with fixed random states. run_comparison(X, y) evaluates all six under a stratified 7:3 split with 5-fold CV, returning accuracy, macro precision, recall and F1, and CV accuracy; run_ablation(X, y) does leave-one-feature-out and single-feature analysis anchored on the random forest (Table 3, B6) to isolate the contribution of each feature descriptor. Both export CSV and LaTeX, run_comparison also renders the best model's confusion matrix, and `python examples/run_benchmark.py` reproduces the full benchmark.

Validity guarantees. Every cloud is resampled to exactly n_{points} points, preventing sampling-density leakage, ensuring volumetric features reflect geometry rather than point-sampling density. All randomness derives from one independent `RandomState` created at construction (not NumPy's global seed), so a given seed yields point-wise identical clouds, features, and tables — asserted by the test suite. There is no compiled code, GPU, or network access at runtime.

4. Illustrative examples

4.1. Quickstart

Install the package and its dependencies first with `pip install -e .` (requires Python $\geq$ 3.8); the API below corresponds to version v0.1.0 and may change in later releases. The following end-to-end example generates a dataset, extracts features, and runs both experiments in eight lines:

```
from solderbench import (SolderJointGenerator, extract_features,
get_baselines, run_comparison, run_ablation)
# generate 1250 joints (5 classes x 250) at seed 42
gen = SolderJointGenerator(n_points=1500, noise=0.02, seed=42)
clouds, labels = gen.generate(n_per_class=250)   # 5 classes x 250 joints
X, y = gen.to_features(clouds, labels)           # 7-feature matrix + labels
results = run_comparison(X, y)   # 6 methods x Acc/P/R/F1 + 5-fold CV
print(results[["no", "name", "Accuracy", "F1"]])
abl = run_ablation(X, y)         # leave-one-out + single-feature dF1
print(abl.head(8))
```

4.2. Benchmark comparison

`examples/run_benchmark.py` produces the benchmark snapshot reported in Table 3: 1250 samples (5 classes $\times$ 250 joints, 1500 points per cloud, seed 42, stratified 7:3 split; 5-fold cross-validation on the training set). Three observations follow. First, the industrial rule S3DA (B1) is a strong baseline at 0.8960 accuracy, confirming

Table 3: Baseline comparison on the Solderbench benchmark (seed = 42, 1250 samples). Best results in bold.

No.	Method	Accuracy	Precision	Recall	macro-F_1	CV Acc. (5-fold)
B1	S3DA (volume threshold)	0.8960	0.8990	0.8960	0.8964	0.9234 ± 0.0138
B2	Logistic Regression	0.9387	0.9416	0.9387	0.9386	0.9657 ± 0.0198
B3	SVM (RBF)	0.9387	0.9390	0.9387	0.9385	0.9771 ± 0.0177
B4	KNN	0.8933	0.8944	0.8933	0.8921	0.9234 ± 0.0168
B5	GradientBoosting	0.9413	0.9417	0.9413	0.9412	0.9577 ± 0.0138
B6	**Random Forest**	**0.9440**	**0.9452**	**0.9440**	**0.9439**	$\mathbf{0.9680 \pm 0.0160}$

that solder volume alone carries most of the class signal — exactly the physics the generator encodes. Second, all learning-based methods improve on the rule, with the random forest best on the held-out test set (accuracy 0.9440, macro-F_1 0.9439). Third, test-set and cross-validation rankings do not fully coincide: the RBF SVM has the highest CV accuracy (0.9771) yet only 0.9387 on the held-out test set, a textbook signature of mild overfitting to the training folds. This is exactly why Solderbench reports both protocols and why users must trust the independent test set rather than cross-validation alone. Figure 1 shows the confusion matrix of the random forest on the held-out test set (375 samples, 21 misclassified). Bridging is perfectly separated (75/75), because the `min_dist` feature exposes the

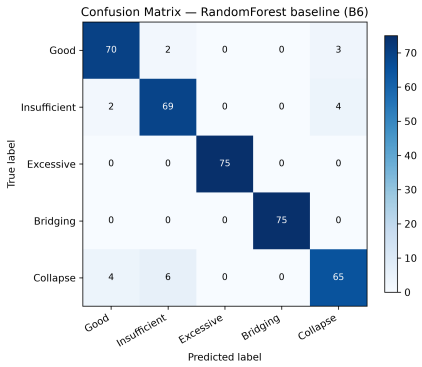


Figure 1: Confusion matrix for the Random Forest (B6) on the held-out test set (375 samples, seed = 42). Bridging is perfectly classified via the `min_dist` feature; the main confusion occurs between collapse and good and insufficient classes.

abnormally small inter-joint distance by construction. All 21 errors stay within the {good, insufficient, collapse} triad: collapse loses 4 samples to good and 6 to insufficient, and gains 3 good samples misclassified as collapse; good and insufficient exchange 2 samples each way, and insufficient loses a further 4 samples to collapse. The overall error therefore remains between collapse and the good and insufficient classes, because a shallow dent alters the height distribution only subtly — precisely why `h_var` — the most informative descriptor in the ablation (Table 4) — is needed to separate it; this mirrors the hard case reported for real 3D inspection systems [6, 7].

4.3. Feature ablation

Table 4 reports the leave-one-feature-out (LOO) ablation for the random forest (full-model macro-F_1 = 0.9439). The ablation has a clean physical reading. Height variance (`h_var`) is the single most valuable descriptor because

Table 4: Leave-one-feature-out ablation of the RandomForest model (macro-F_1; full model = 0.9439).

Removed feature	ΔF_1 **vs. full-model** F_1	**Interpretation**
h_var	-0.0459	Most informative: collapse dents and excessive bumps reshape the height distribution
vol_area_ratio	-0.0105	Compactness cue separating volumetric defects from nominal joints
min_dist	-0.0054	Small overall cost because bridging is nearly separable already
area	-0.0052	Correlated with vol; adds modest independent signal
convex_vol_ratio	-0.0025	Irregularity cue; partially redundant with h_var
vol	+0.0026	Slightly redundant once ratios are present
max_h	0.0000	Fully redundant given h_var and vol

the two hardest defects — collapse (a dent) and excessive solder (a bump) — are primarily height-distribution phenomena, whereas volume-only cues saturate. Single-feature runs corroborate this: the strongest single feature is `convex_vol_ratio` (macro-$F_1 = 0.6412$), which captures shape irregularity common to several defect types, while the weakest is `min_dist` (macro-$F_1 = 0.3616$), which is highly informative for bridging but carries no signal for the other four classes. The near-zero or positive deltas for `vol` and `max_h` indicate information redundancy rather than irrelevance: both are strongly correlated with the ratio features. Volume alone is in fact a strong single cue — it underpins the industrial S3DA rule (B1), which reaches 0.8960 accuracy using volume only — and its near-zero delta here reflects redundancy with the ratio features, not uselessness. Removing `vol` slightly increases macro-F_1 within small numerical fluctuations; this is purely a statistical artifact given that volume information is already encoded within ratio-based descriptors.

Table 5: Selected representative single-feature macro-F_1 scores (RandomForest, seed = 42). Full results are programmatically exported by run_ablation.

Single feature	**macro-F_1**
`convex_vol_ratio`	0.6412
`min_dist`	0.3616

5. Impact

Solderbench contributes to three communities. For inspection-algorithm research, it is the first public, reproducible 3D benchmark for solder-joint defect classification, replacing per-paper private datasets with a common task, protocol, and versioned reference numbers; new methods are ranked against six baselines with one command. For teaching, it is a self-contained geometric-ML case study that fits a short lab session and runs on any laptop, with each design decision doubling as a benchmark-methodology lesson. For reproducibility, seed-level determinism, a one-command script, machine-readable exports, a `codemeta.json` descriptor, and an MIT license make the software trivially citable, archivable, and reusable. The path to production data is short: a user with a real 3D AOI camera replaces the body of `load_dataset` — or passes their own X, y — with code that reads scans, segments joints, and calls `extract_features`; all downstream experiments run unchanged. Beyond the classical-ML workflow, the generator's raw point clouds can be passed directly to external deep models (e.g., PointNet [8, 9]); Solderbench supports both scenarios.

6. Limitations

We state the limitations of Solderbench honestly. The generator is a geometric model, not a physics simulator: it does not model wetting dynamics, surface tension, oxidation, or realistic sensor artifacts (occlusion, multi-path reflections, structured-light failure), and its noise is a simple isotropic perturbation; class distributions are plausible and overlapping, not fitted to production data. Three scope limits remain. First, the five classes are idealized single-fault geometries — compound defects (e.g., insufficient and collapsed together) are unsupported, although they occur on real-world production lines. Second, the clouds model only the individual joint, omitting substrate, pads, and pins, so board-level context is out of scope. Third, resampling every cloud to a fixed n_{points} simplifies real sensors, which return a variable point-count, and may alter the point-density distribution. **Consequently, Solderbench is suitable for algorithm prototyping, fair cross-method comparison, and descriptor-level analysis, but its metric values must not be read as predictions of real industrial AOI performance.** Synthetic validation cannot substitute for validation on real scans; the benchmark guarantees only that comparisons made on it are fair, controlled, and exactly reproducible — a much-needed property currently missing from this research field.

Data availability. The Solderbench source code and all reproduction scripts are released under the MIT license and archived on Zenodo (doi:`10.5281/zenodo.XXXXXXX`); the permanent GitHub repository is `https://github.com/OWNER/solderbench`. The benchmark dataset is not stored as a fixed file but generated deterministically at runtime by running `python examples/run_benchmark.py`, which regenerates every table and figure in this paper.

7. Conclusions

We presented Solderbench, an open-source, pure-Python synthetic benchmark for 3D PCB solder-joint defect inspection with a parametric five-class generator, seed-level reproducibility, a seven-dimensional feature extractor, six baselines including the industrial S3DA rule, and automated comparison and ablation runners. On the reference benchmark (1250 samples, seed 42) a random forest reaches 0.9440 accuracy and 0.9439 macro-F_1. By removing the data-access barrier that fragments 3D solder-joint inspection, Solderbench makes algorithm development, comparison, and teaching open and reproducible. Future work: (i) fitting generator parameters to measured scans; (ii) adding sensor-artifact models; (iii) compound-defect taxonomy; (iv) real-data conversion scripts; (v) deep point-cloud model interfaces (e.g. PointNet, PointNet++); and (vi) preset difficulty templates.

8. Acknowledgements

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors. Computations were performed with the Solderbench toolkit (MIT license).

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

CRediT authorship contribution statement

Qinmin Wang: [TODO: CRediT roles] **Zixin Wang:** [TODO: CRediT roles] **Hongyuan Wang:** [TODO: CRediT roles] **Zongzhong Wang:** [TODO: CRediT roles — corresponding author]

References

[1] Q. Ling, N.A.M. Isa, Printed circuit board defect detection methods based on image processing, machine learning and deep learning: A survey, IEEE Access 11 (2023) 15921-15944. https://doi.org/10.1109/ACCESS.2023.3245093

[2] Y. Zhou, M. Yuan, J. Zhang, G. Ding, S. Qin, Review of vision-based defect detection research and its perspectives for printed circuit board, Journal of Manufacturing Systems 70 (2023) 557-578. https://doi.org/10.1016/j.jmsy.2023.08.019

[3] M. Moganti, F. Ercal, Automatic PCB inspection systems, IEEE Potentials 14(3) (1995) 6-10. https://doi.org/10.1109/45.464686

[4] W. Dai, A. Mujeeb, M. Erdt, A. Sourin, Soldering defect detection in automatic optical inspection, Advanced Engineering Informatics 43 (2020) 101004. https://doi.org/10.1016/j.aei.2019.101004

[5] S. Tang, F. He, X. Huang, J. Yang, Online PCB defect detector on a new PCB defect dataset, arXiv preprint arXiv:1902.06197 (2019). https://doi.org/10.48550/arXiv.1902.06197

[6] Y. Xu, Z. Wang, J. Liang, Y. Zhang, J. Feng, W. Miao, J. Wang, G. Pu, S3DA: A 3D point cloud based PCB solder defect detection algorithm, in: S. Liu (Ed.), Software Fault Prevention, Verification, and Validation (SFPVV 2024), Lecture Notes in Computer Science, vol. 15393, Springer, 2025, pp. 115-130. https://doi.org/10.1007/978-981-96-1621-3_8

[7] Q. Hu, K. Hao, B. Wei, H. Li, An efficient solder joint defects method for 3D point clouds with double-flow region attention network, Advanced Engineering Informatics 52 (2022) 101608. https://doi.org/10.1016/j.aei.2022.101608

[8] C.R. Qi, H. Su, K. Mo, L.J. Guibas, PointNet: Deep learning on point sets for 3D classification and segmentation, in: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2017, pp. 652-660. https://doi.org/10.1109/CVPR.2017.16

[9] C.R. Qi, L. Yi, H. Su, L.J. Guibas, PointNet++: Deep hierarchical feature learning on point sets in a metric space, in: Advances in Neural Information Processing Systems 30 (NeurIPS), 2017, pp. 5099-5108. https://proceedings.neurips.cc/paper/2017/hash/d8bf84be3800d12f74d8b05e9b89836f-Abstract.html

[10] Q.-Y. Zhou, J. Park, V. Koltun, Open3D: A modern library for 3D data processing, arXiv preprint arXiv:1801.09847 (2018). https://arxiv.org/abs/1801.09847

[11] L. Breiman, Random forests, Machine Learning 45(1) (2001) 5-32. https://doi.org/10.1023/A:1010933404324

[12] C. Cortes, V. Vapnik, Support-vector networks, Machine Learning 20(3) (1995) 273-297. https://doi.org/10.1007/BF00994018

[13] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, E. Duchesnay, Scikit-learn: Machine learning in Python, Journal of Machine Learning Research 12 (2011) 2825-2830.

[14] C.R. Harris, K.J. Millman, S.J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, E. Wieser, J. Taylor, S. Berg, N.J. Smith, et al., Array programming with NumPy, Nature 585 (2020) 357-362. https://doi.org/10.1038/s41586-020-2649-2

[15] P. Virtanen, R. Gommers, T.E. Oliphant, M. Haberland, T. Reddy, D. Cournapeau, E. Burovski, P. Peterson, W. Weckesser, J. Bright, et al., SciPy 1.0: Fundamental algorithms for scientific computing in Python, Nature Methods 17 (2020) 261-272. https://doi.org/10.1038/s41592-019-0686-2

[16] W. McKinney, Data structures for statistical computing in Python, in: Proceedings of the 9th Python in Science Conference, 2010, pp. 51-56. https://doi.org/10.25080/Majora-92bf1922-00a

[17] J.D. Hunter, Matplotlib: A 2D graphics environment, Computing in Science & Engineering 9(3) (2007) 90-95. https://doi.org/10.1109/MCSE.2007.55